

Forecasting the Unexpected

Anomaly Detection Without a Crystal Ball (Just Spikes)

Developed for Reply and brought to you by the team:

NAPOLETANI A ROCCARASO

Gabriele Rizzo

Annalaura Granata

Daniele Lupico



Introduction to Our Work

Time Series Analysis & Anomaly Detection

Project Objective

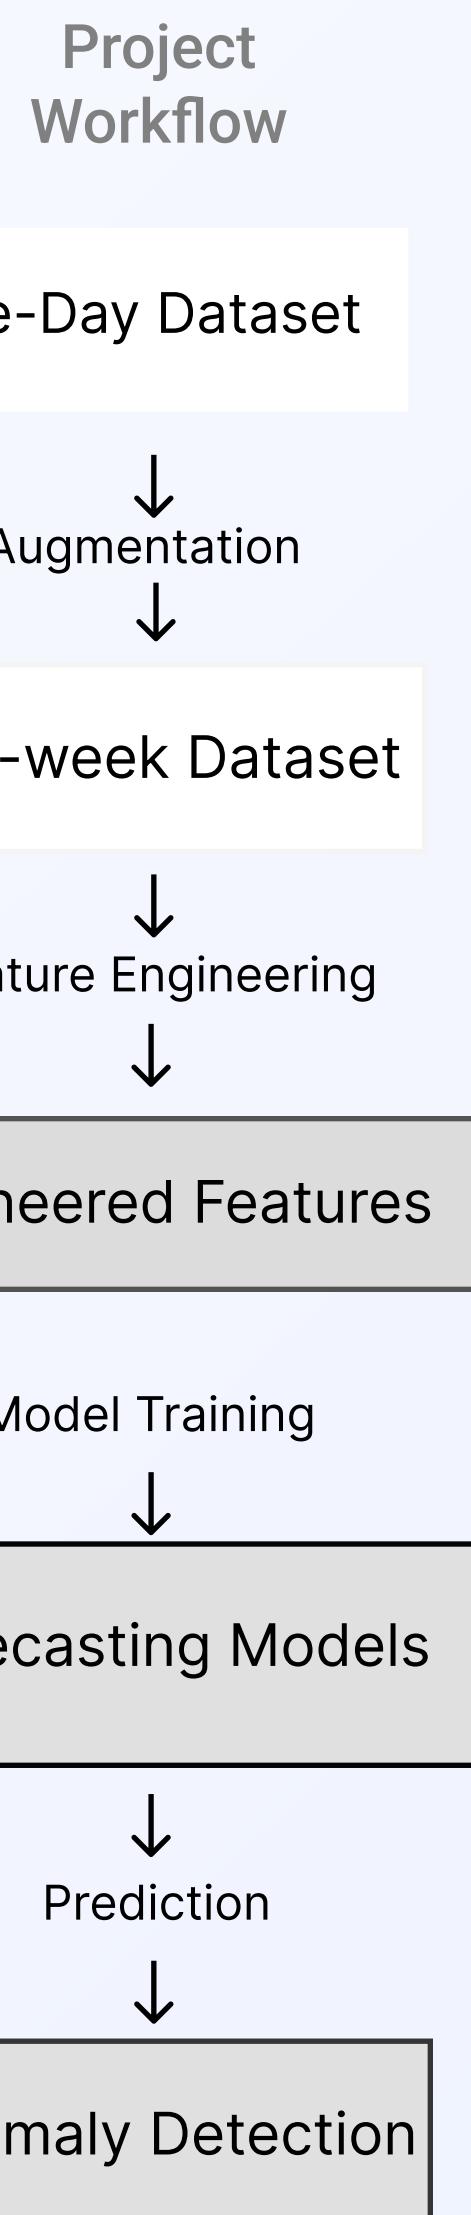
- ✓ Transform a one-day transaction dataset into a full week of **realistic data**
- ✗ Engineer useful features for improved analysis
- Apply **forecasting** and **anomaly detection** models

Implementation Tools

seaborn

scikit-learn

tensorflow



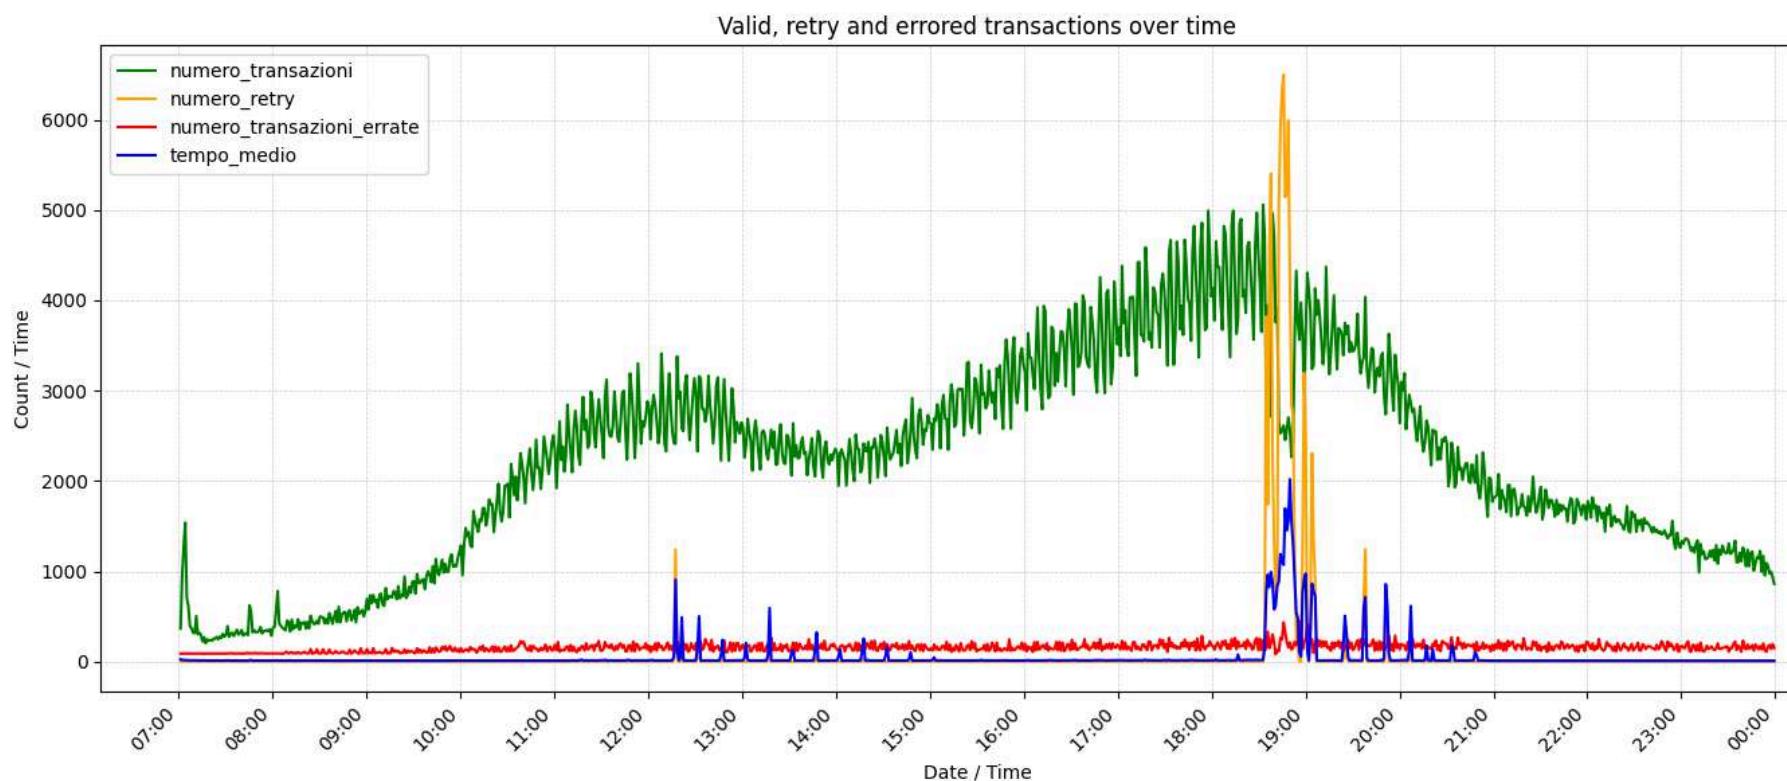
Dataset and Augmentation Strategy

Extending a single-day transaction dataset to a full week of realistic data

Original Dataset

From Client-provided fields, we selected the following:

- ↔ numero_transazioni
- ⌚ tempo_medio
- ⌚ numero_retry
- ⚠️ numero_transazioni_errate



💡 Augmentation Strategy



Predefined Rules

- ✓ Augmentation with clipped noise
- ✓ Augmentation smoothed outliers, avoiding the model to be too influenced by few extreme spikes
- ✓ Pattern preservation

Advanced Techniques

- ✓ Rule-based noise injection based on variable probability created labeled test scenarios
- ✓ "Slow Day" introduction helped testing model robustness

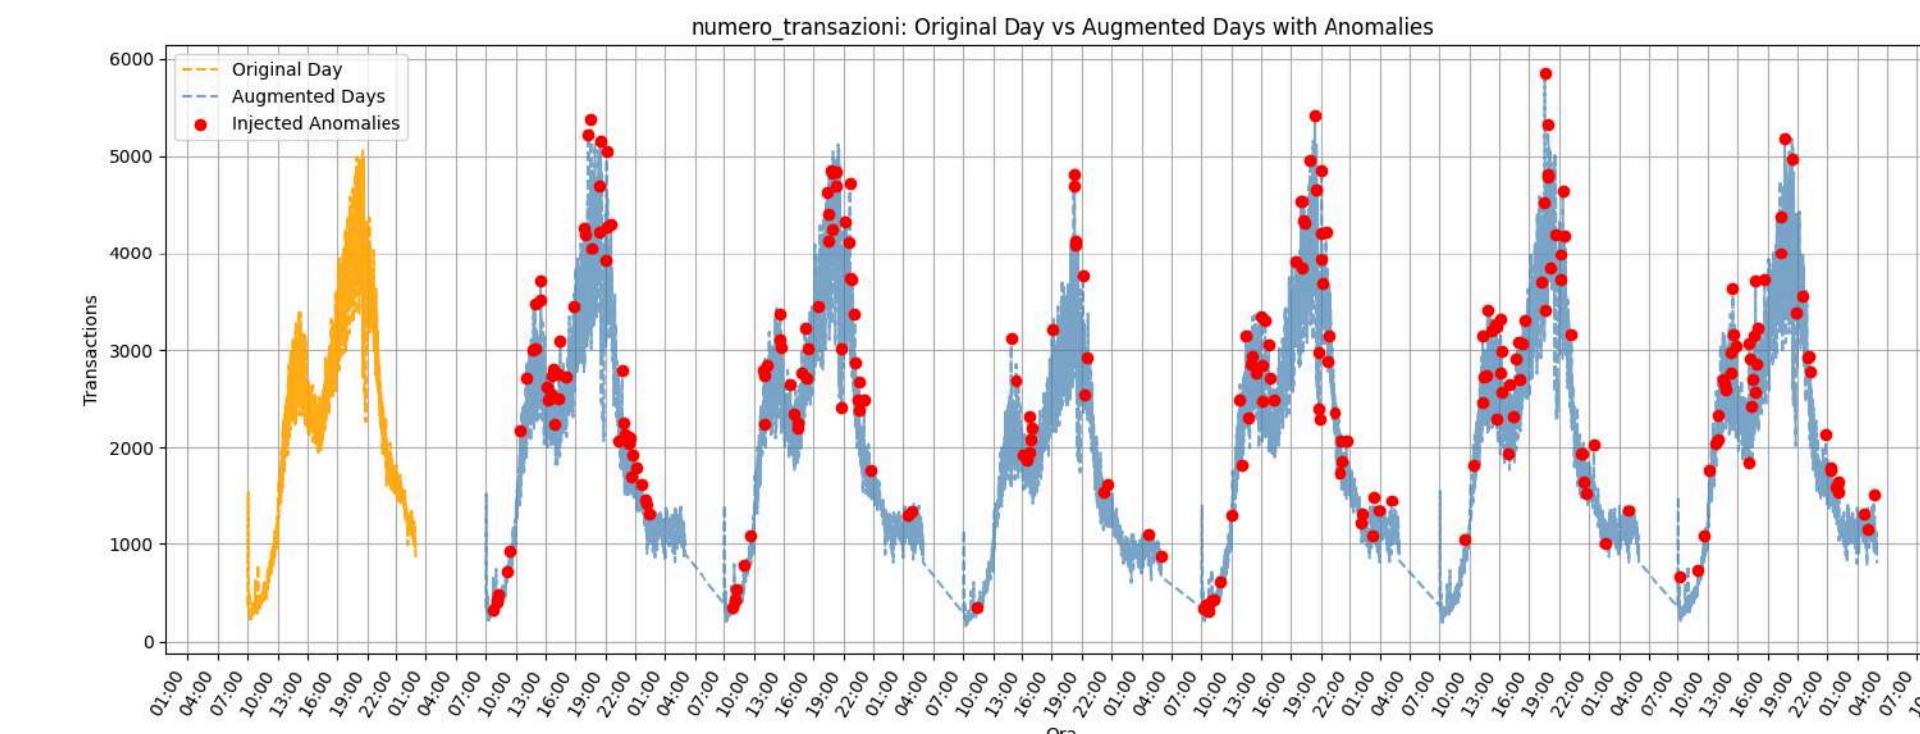
Data Integrity Preservation

Correlation Preservation

- ⌚ Maintained relationships between variables
- ⌚ Realistic dependency structures

Rhythm Preservation

- ⌚ Business cycle simulation: handled interval 00:00-02:59 and represented night time business pause 03:00-06:59



Exploratory Data Analysis

Visualizing and comparing distributions between original and generated datasets

Distribution Comparisons

Original (1 Day)

Generated (1 Week)

Verification Points

Controlled Outliers

- tempo_max and tempo_min used for clipping prevents unrealistic values

Coherent Distributions

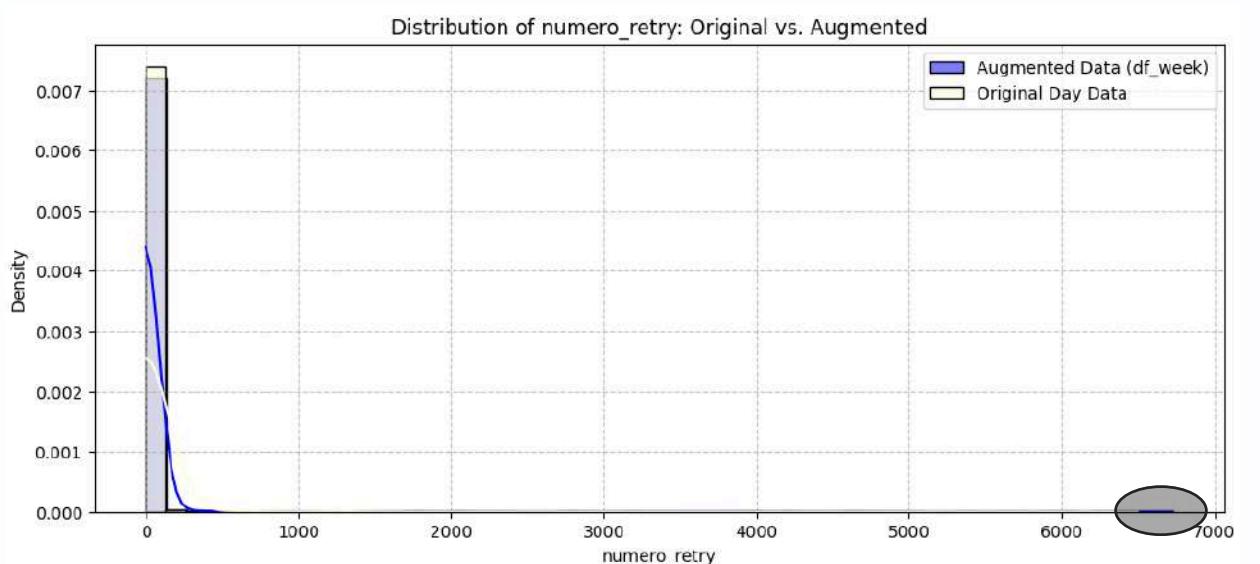
Generated data preserves key statistical properties:

- Similar means and medians
- Comparable variance
- Maintained correlation structure

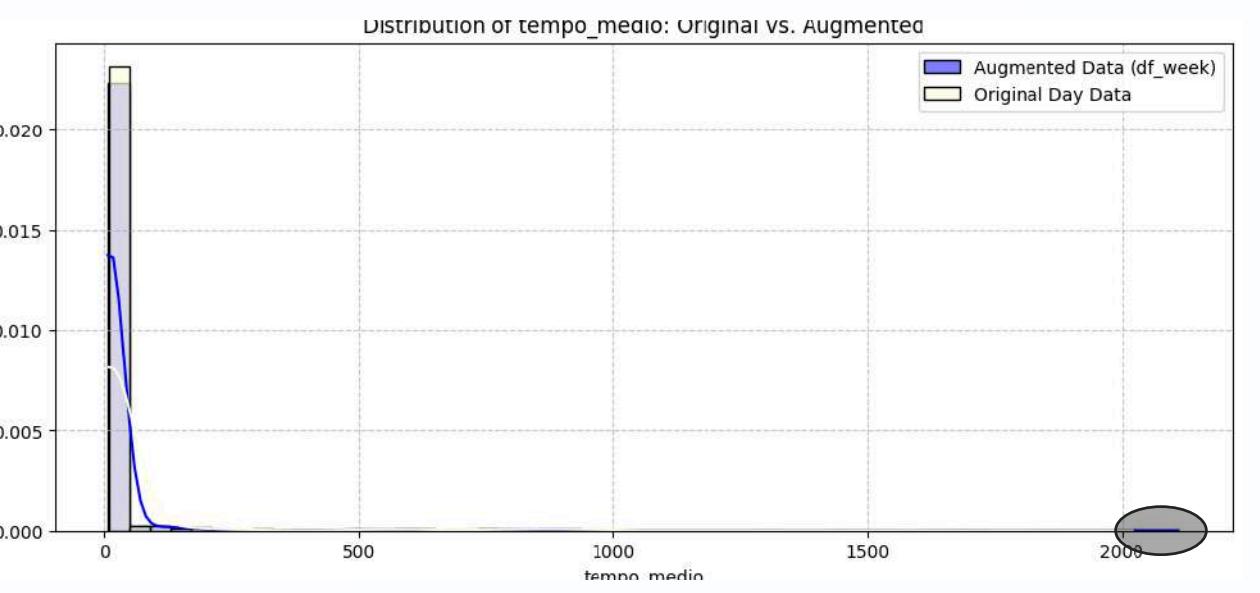
Realistic Noise

- Non-destructive variation
- Natural pattern fluctuations
- Balanced signal/noise ratio

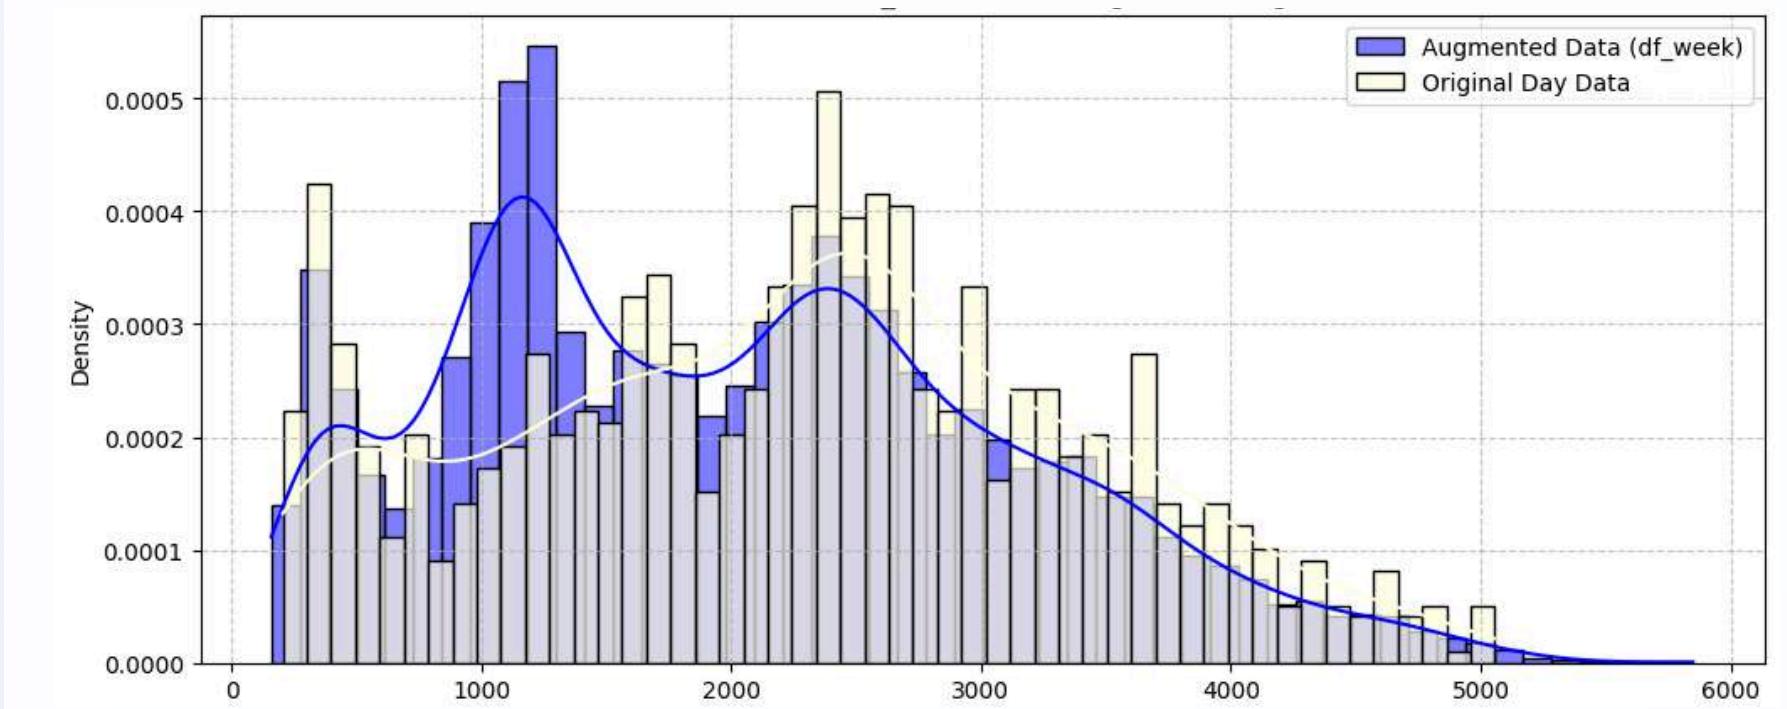
Retry Count



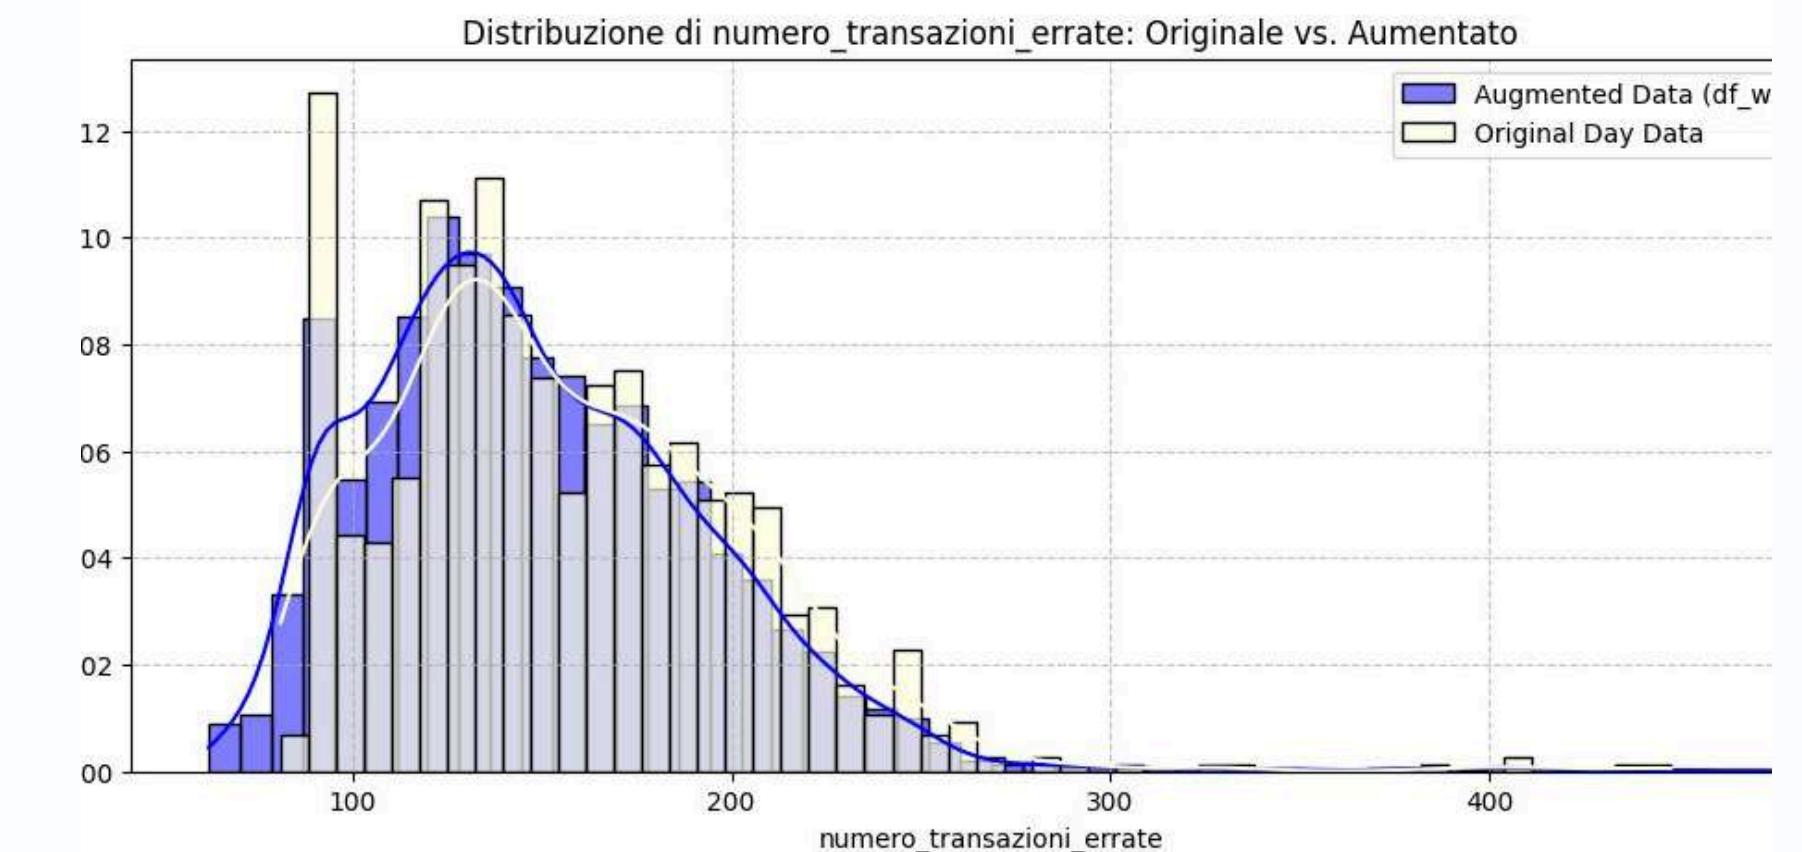
Average Response Time



Number of Transaction

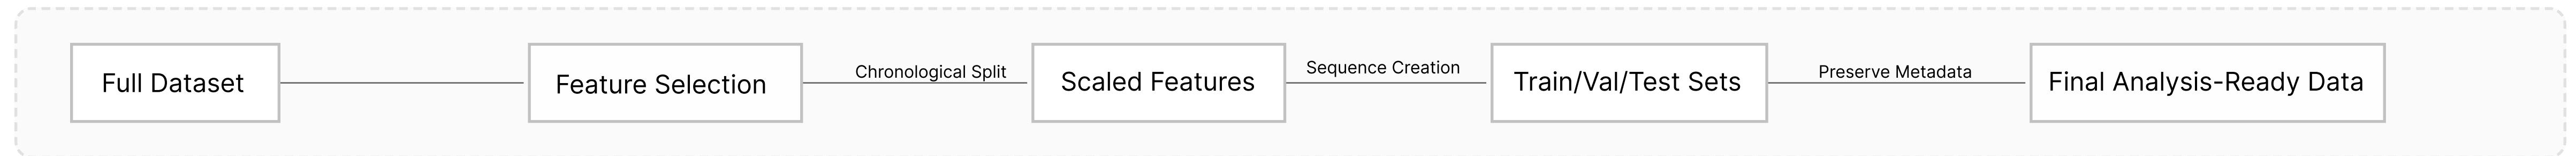


Number Failed Transactions



Final Dataset Preparation

Transforming the engineered dataset into a format ready for modeling



Filter Feature Selection

Removed Features

⌚ tempo_min ⌚ tempo_max

These features were redundant with other engineered metrics and contained outliers that could negatively impact model performance.

Retained Features

numero_transazioni tempo_medio numero_retry
numero_transazioni_errate hour_of_day day_of_week
is_anomaly

- By extracting from data_ora features such as day_of_week and hour_of_day, we empowered our models to learn richer temporal patterns and relationships.
- Data_ora used as index
- In modeling phase only is_anomaly==0 considered

Delta Feature Scaling

StandardScaler applied to numeric features:

- ✓ Zero mean, unit variance transformation
- ✓ Trained on training set only
- ✓ Applied to validation and test sets

Puzzle Data Splitting

Split Proportions

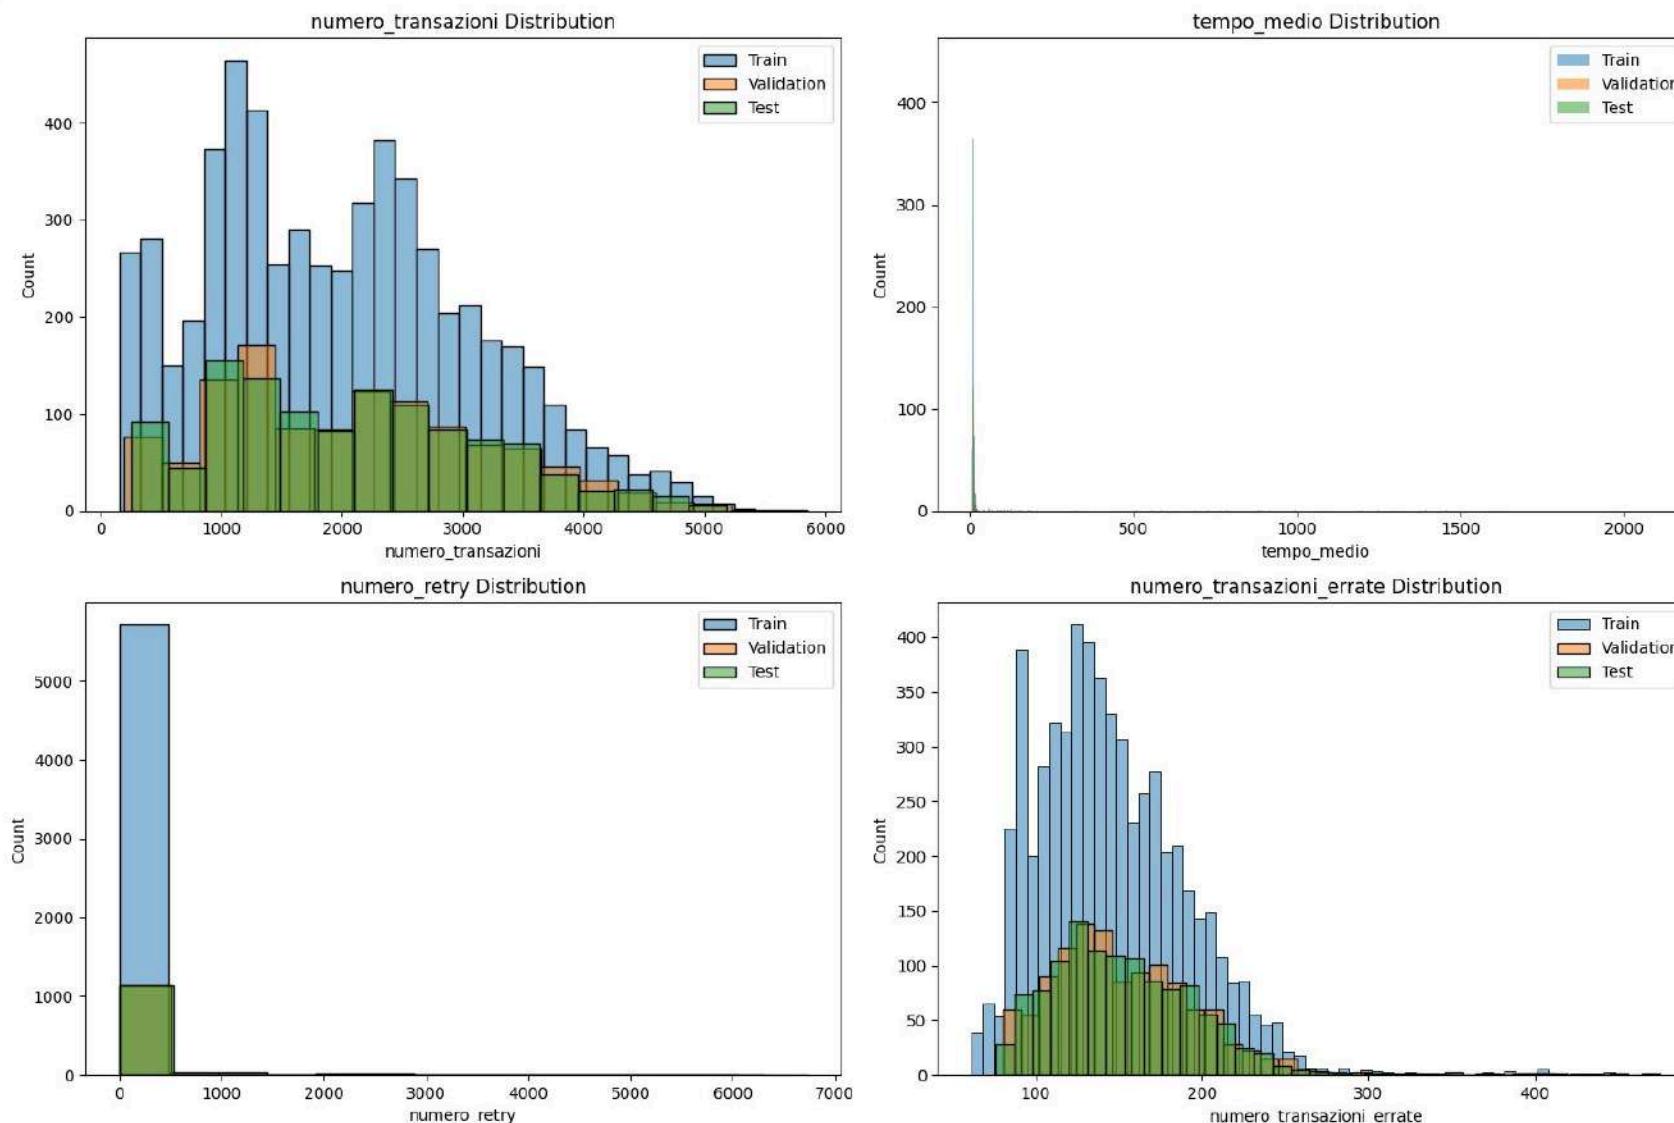
70% Train

15% Val

15% Test

Chronological split to maintain temporal integrity of data

Preserved Metadata



Flag Final Dataset Results

~8200 usable rows (approximately 7 days of data)

✓ Feature-rich dataset ready for time series modeling

Forecasting Models

Predictive approaches used to forecast transaction metrics

⚙️ TCN Model Architecture

↔ Model Selection Approach

Instead of Traditional:

🚫 ARIMA

🚫 Prophet

We Used:

✓ TCN

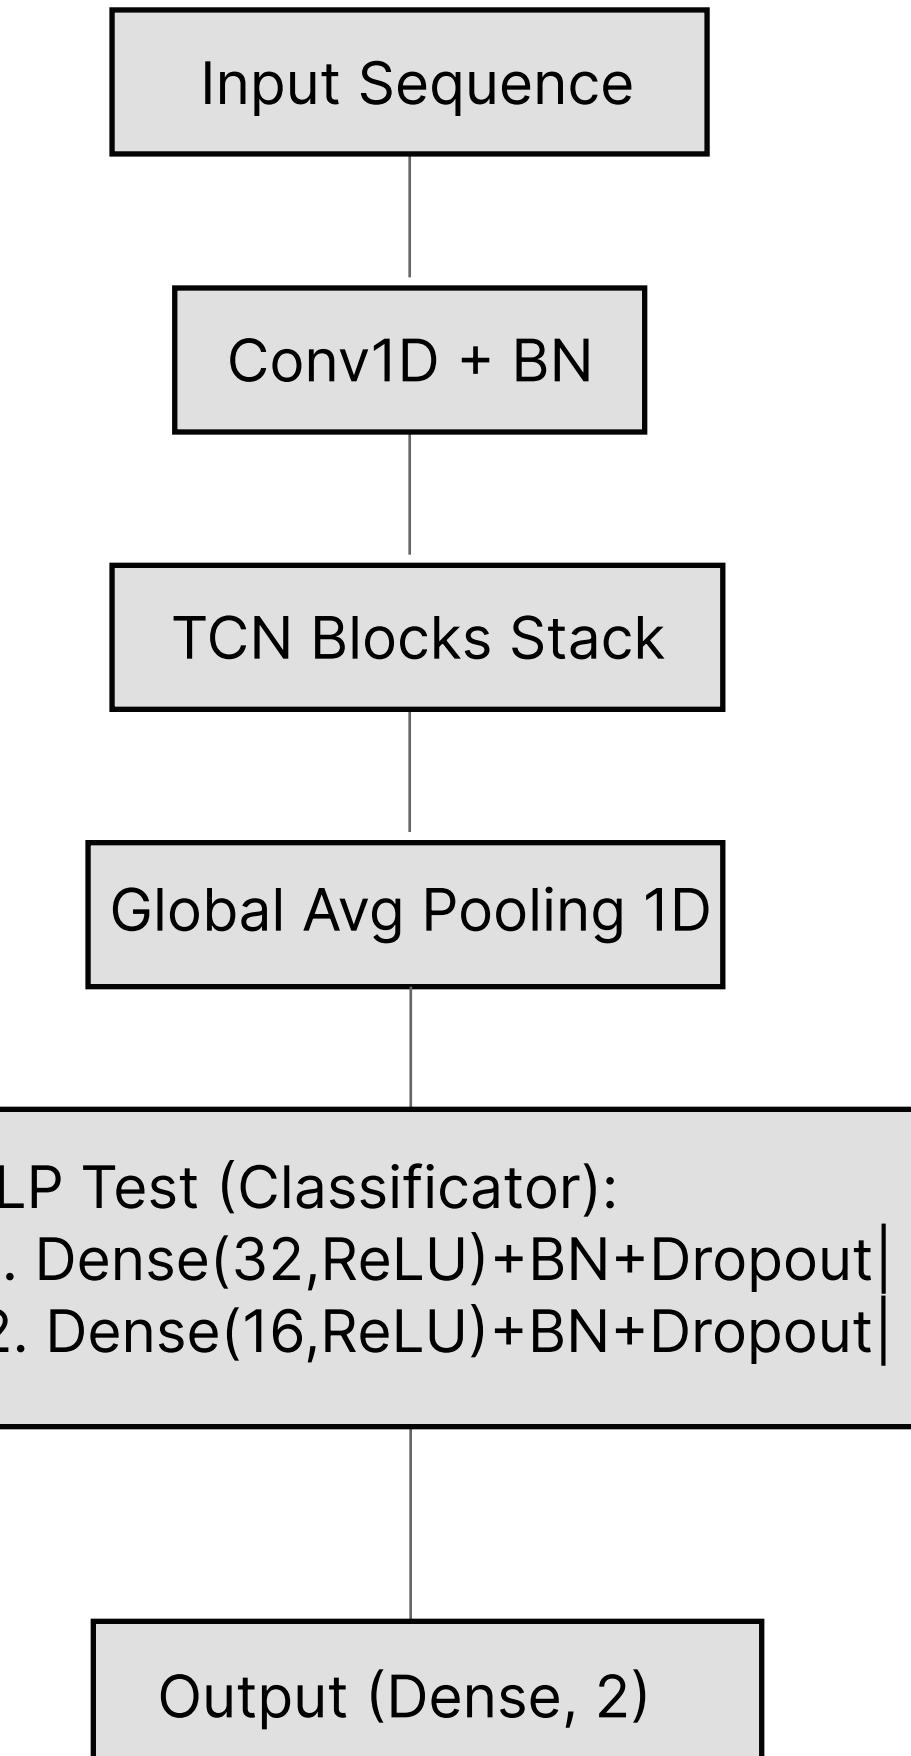
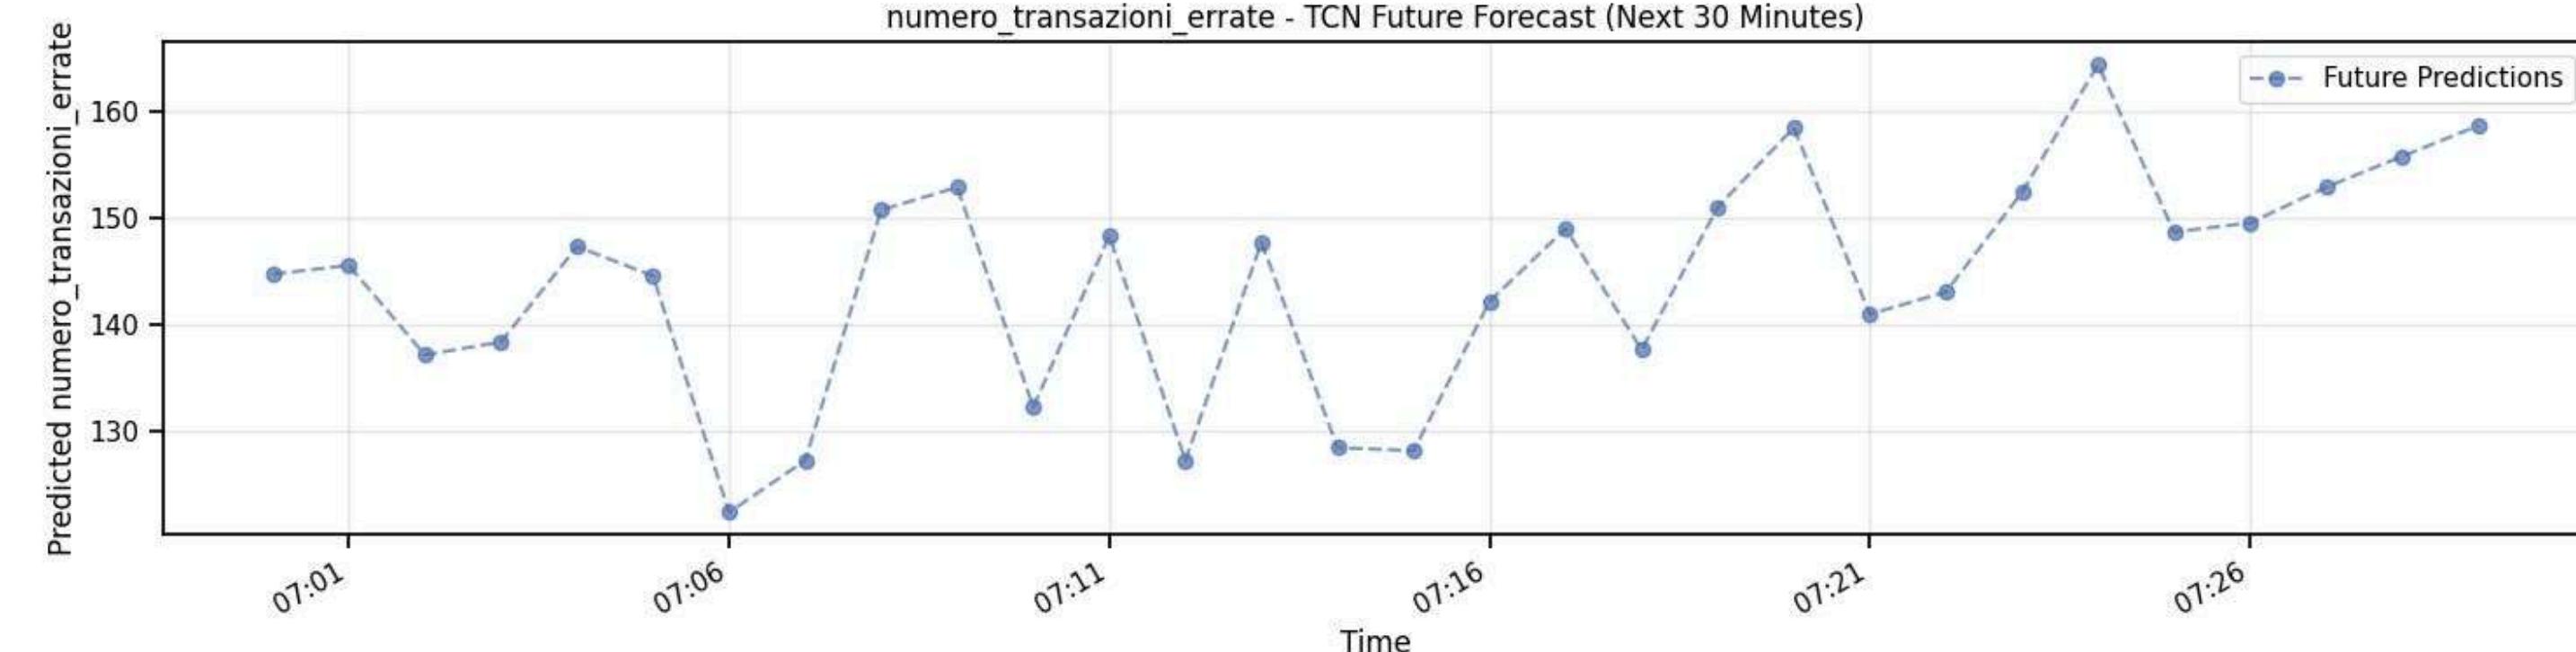
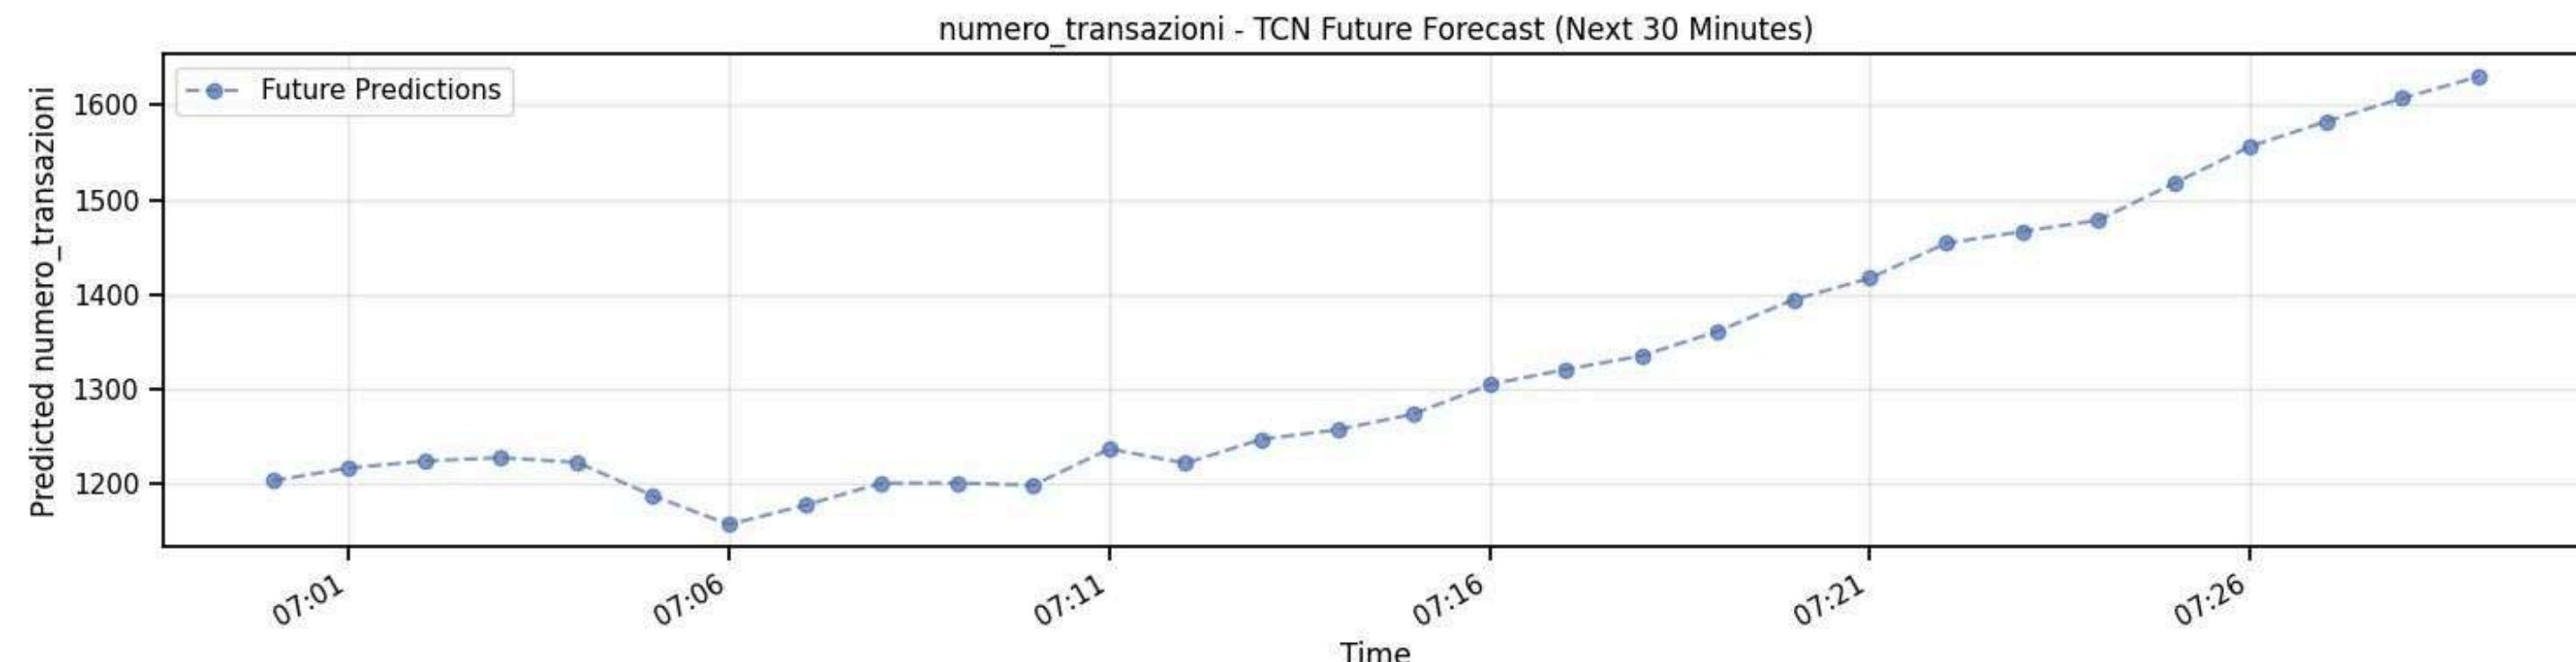
✓ LSTM

Forecasting Goal Features:

- Numero_transazioni
- Numero_transazioni_errate

Key Results:

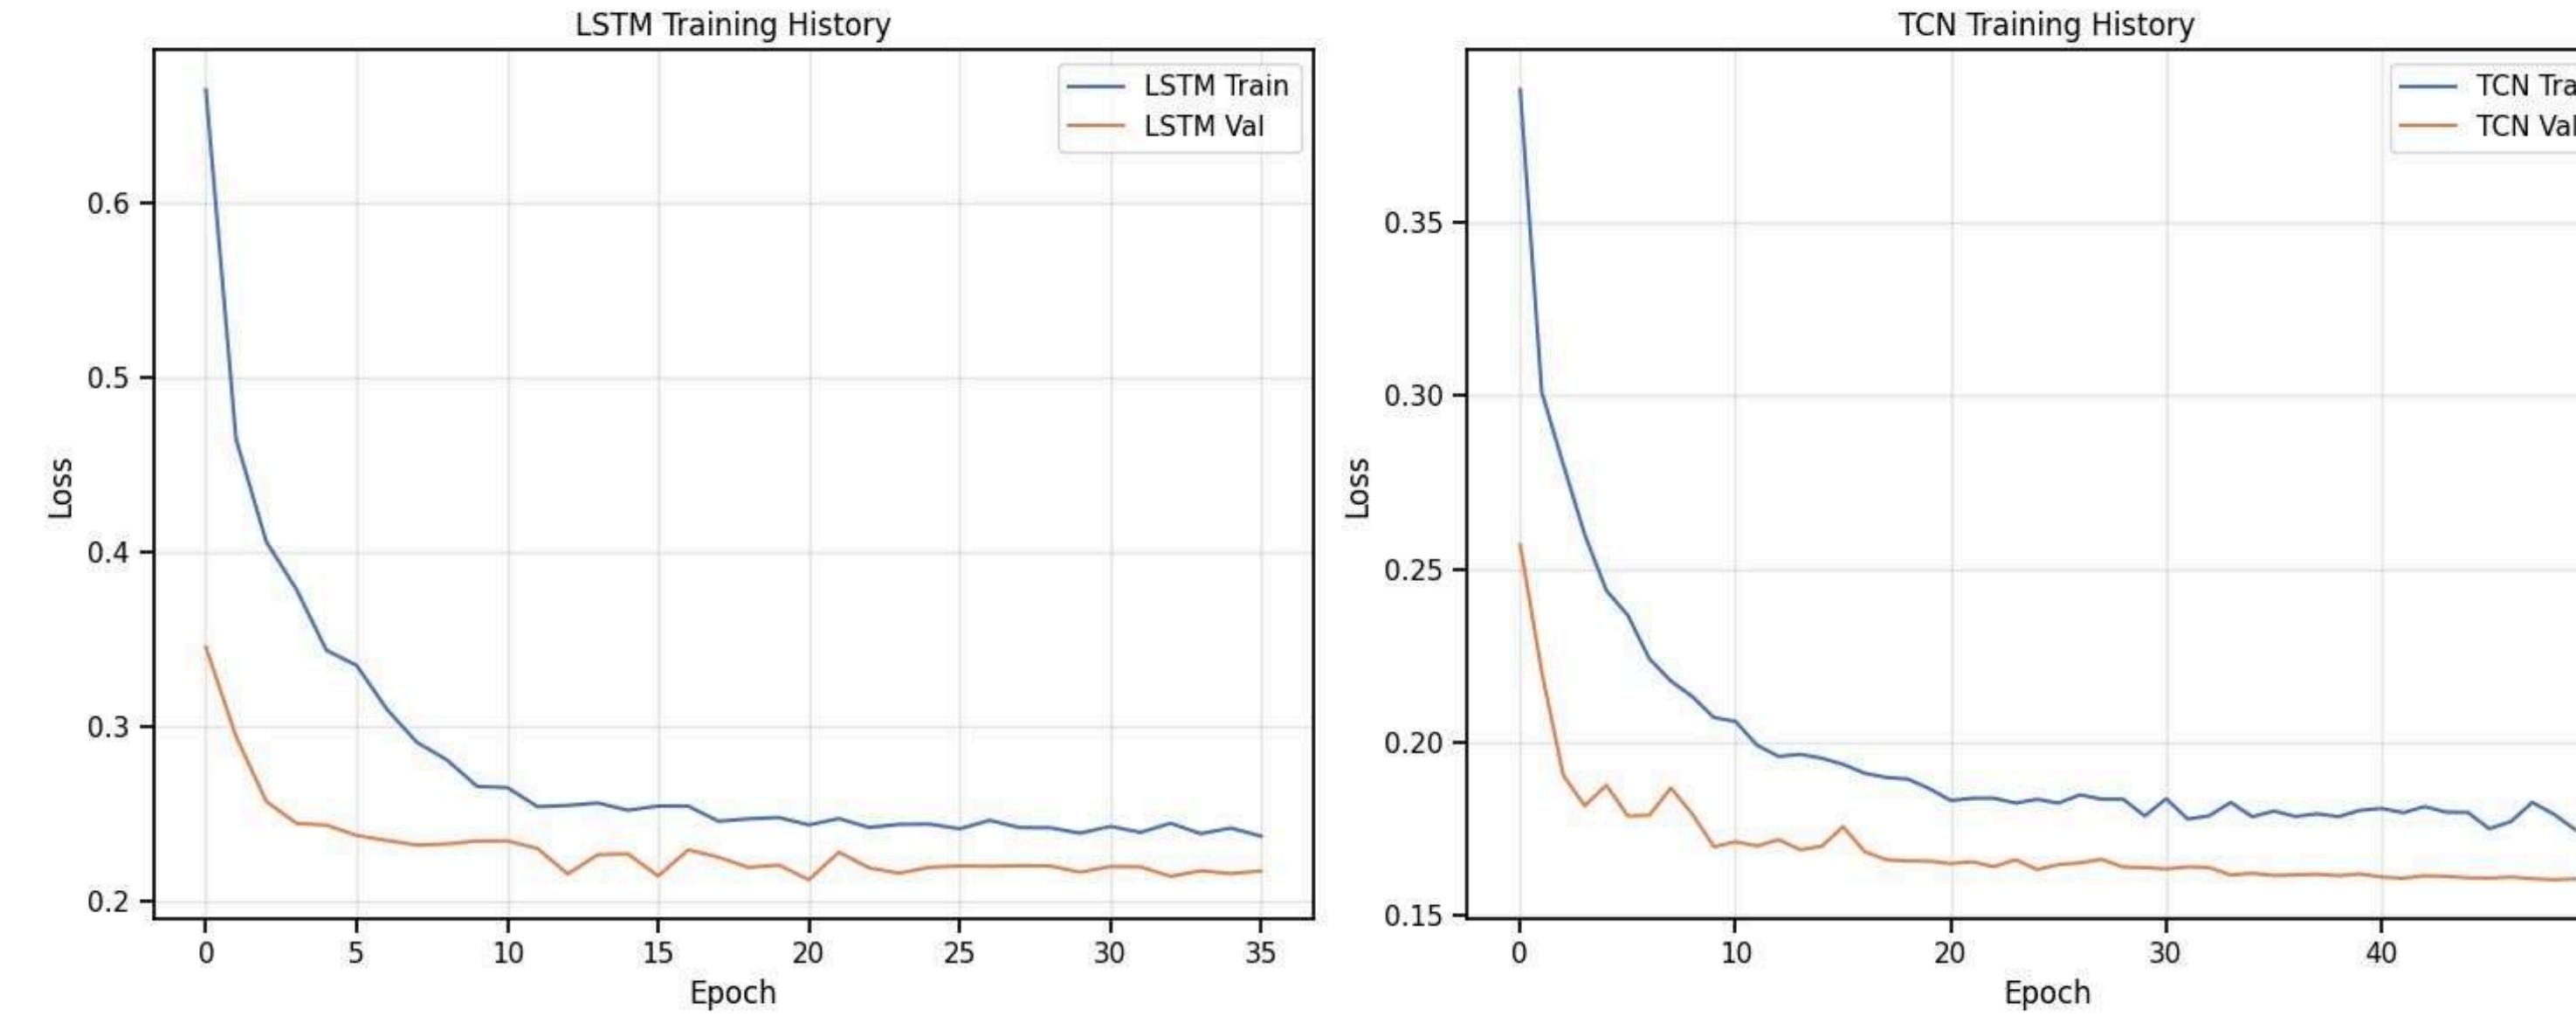
- Predicts 30 minutes ahead
- Utilizes engineered temporal features



Forecast Evaluation

Assessing forecast accuracy and utility for anomaly detection

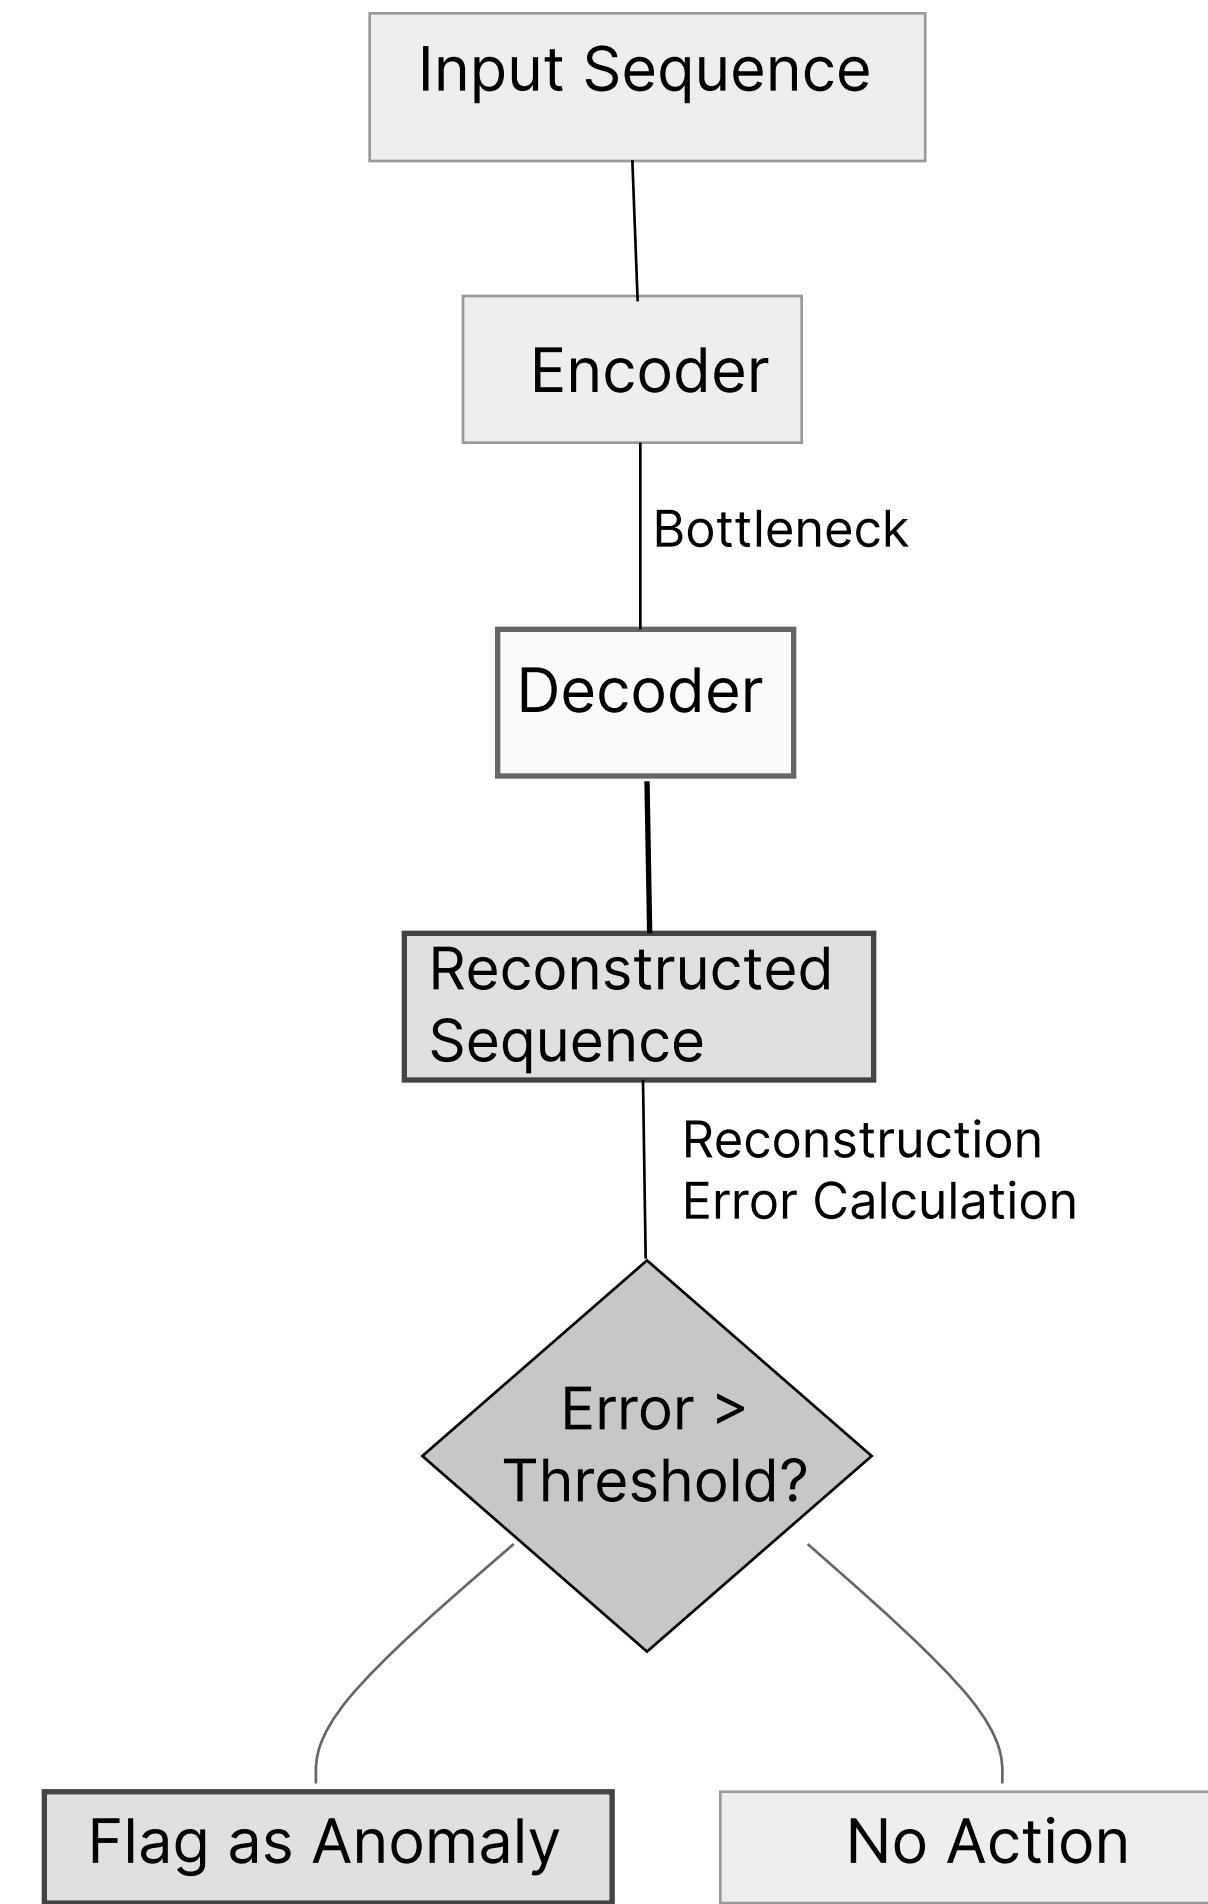
↳ Evaluation Metrics



Anomaly Detection

Using model predictions to identify and flag abnormal system behavior

⚠ Anomaly Detection Approach

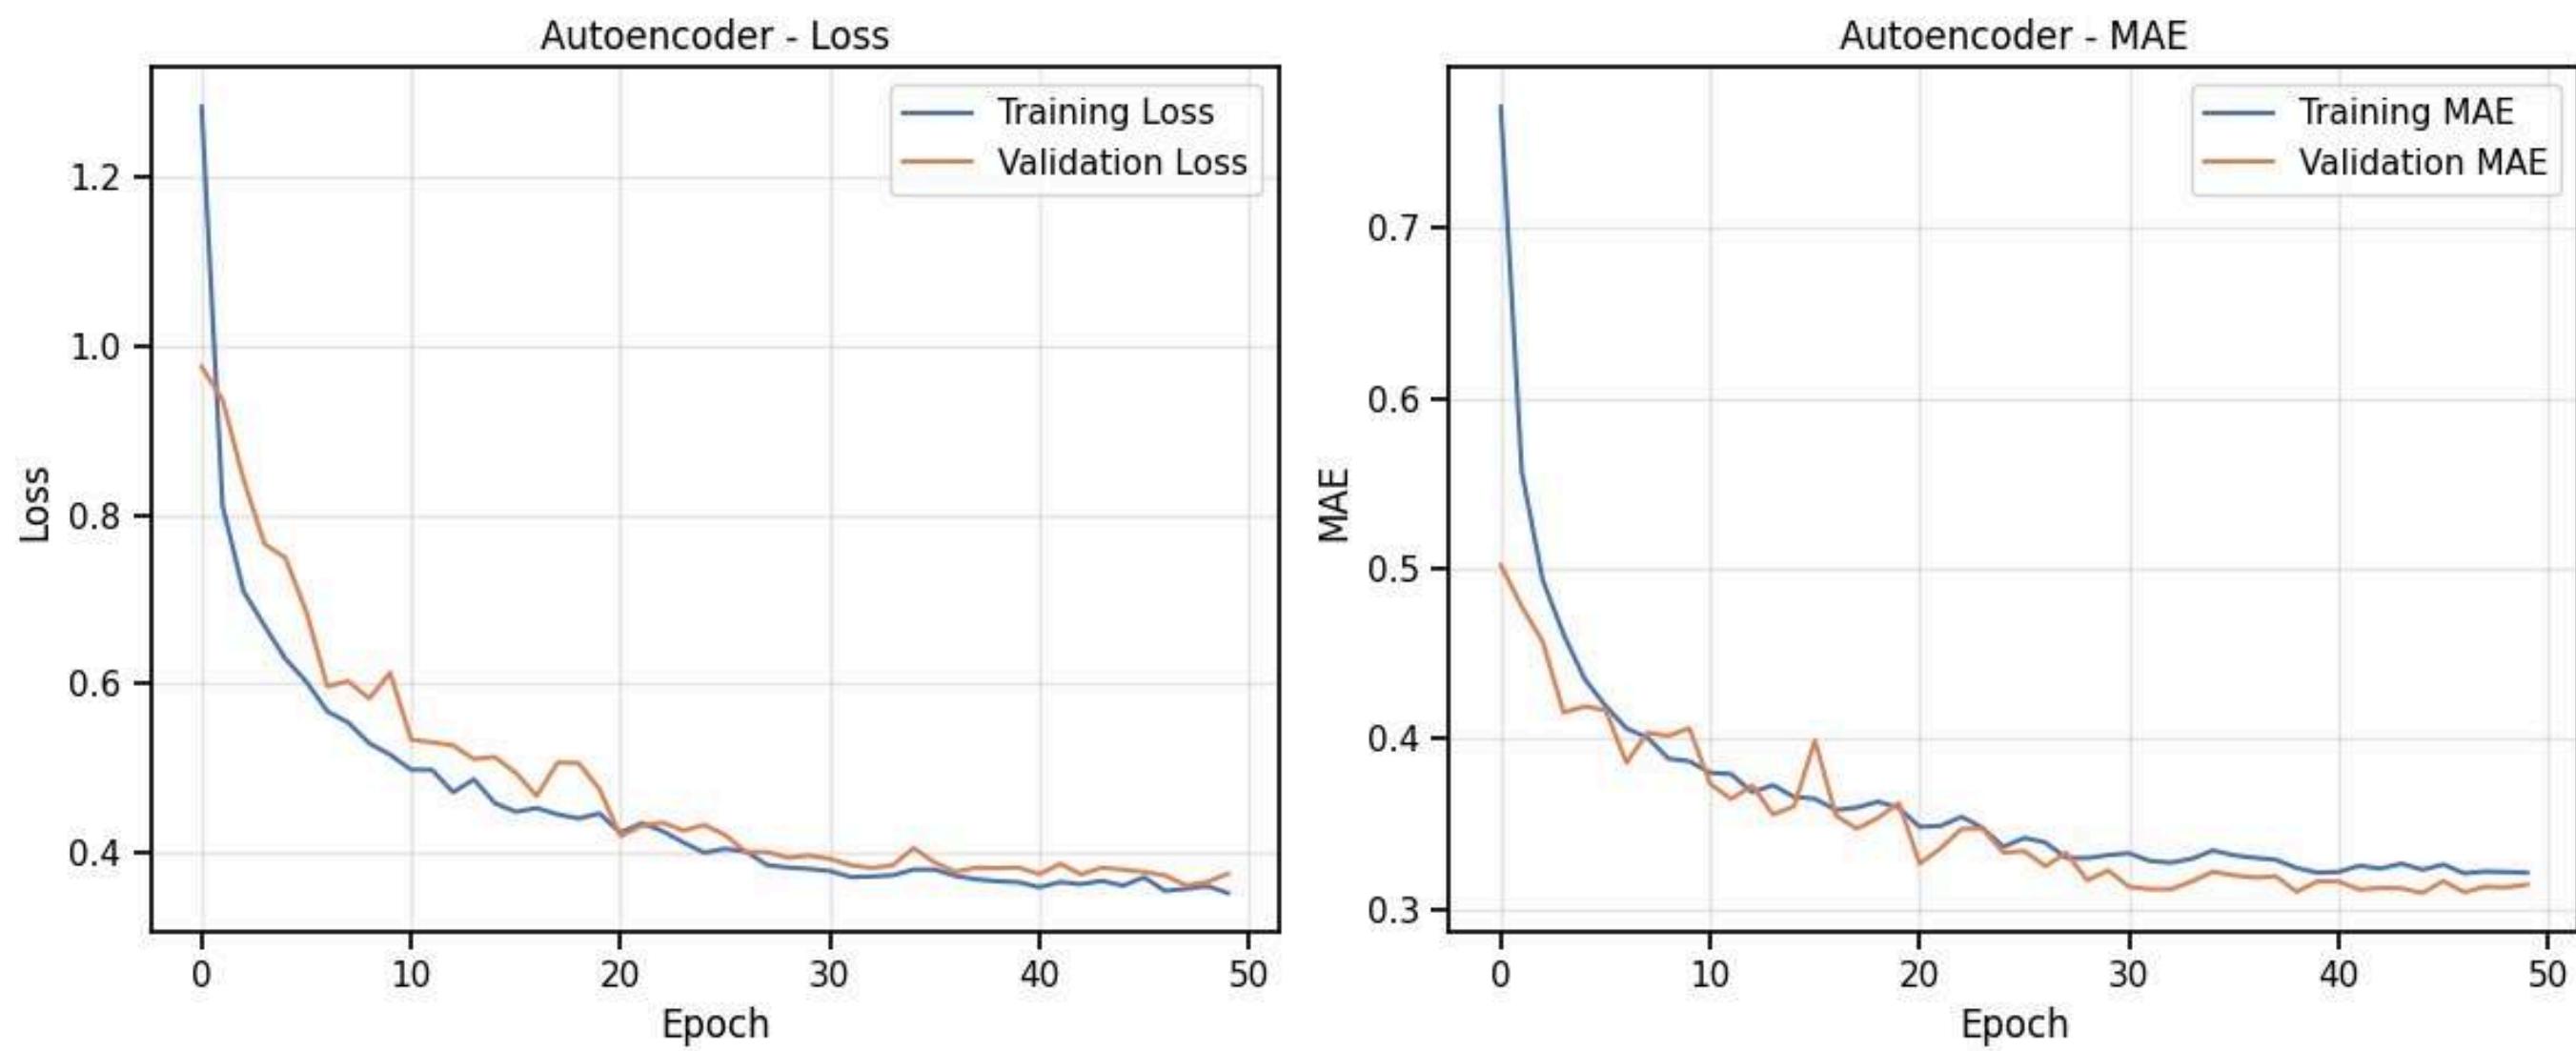


Autoencoder LSTM that detects anomalies based on input reconstruction error

Learns normal patterns

High reconstruction error signals anomaly

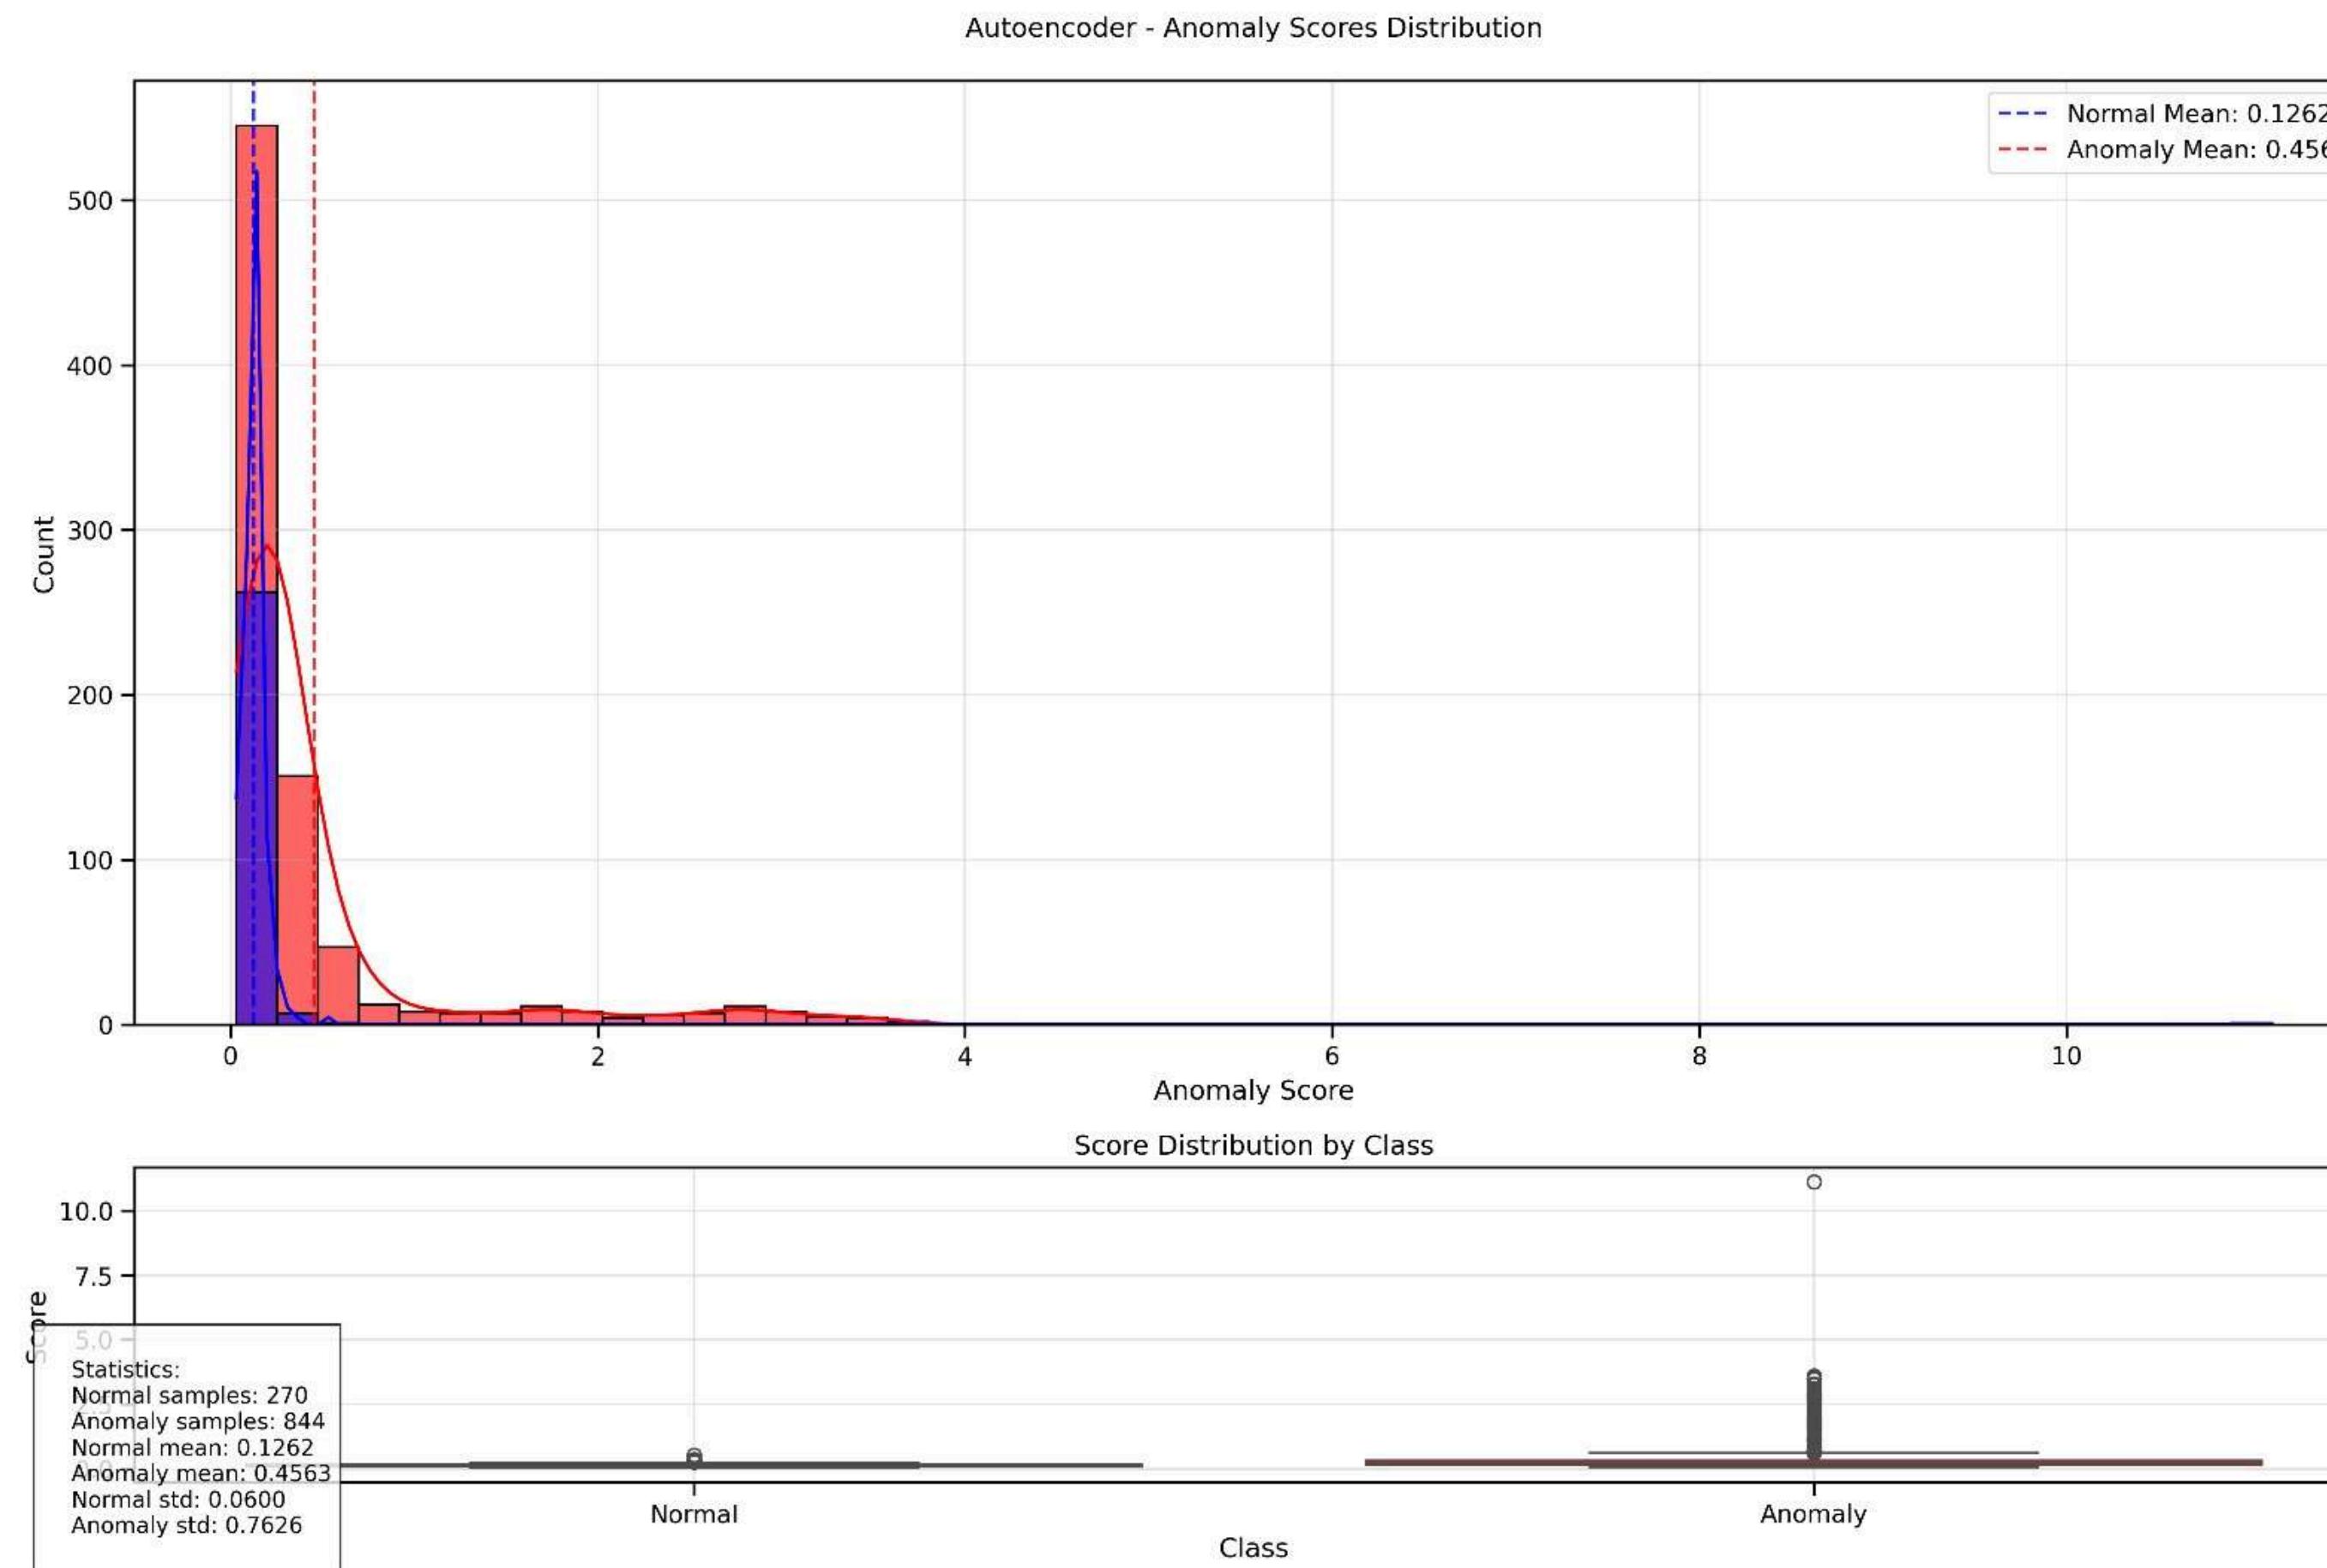
Trained only on normal data



Anomaly Detection

Using model predictions to identify and flag abnormal system behavior

☰ Anomaly Thresholds



Anomaly Detection

Using model predictions to identify and flag abnormal system behavior

Accuracy

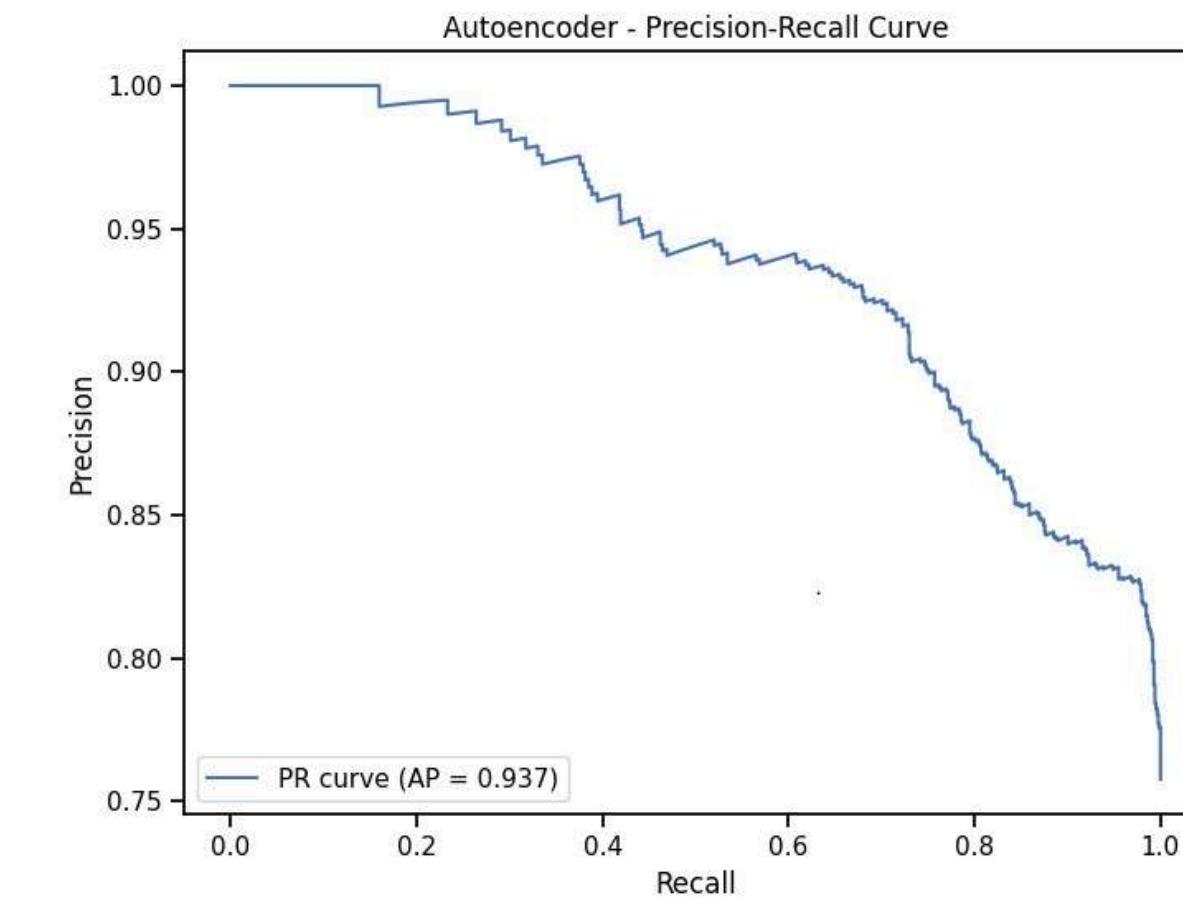
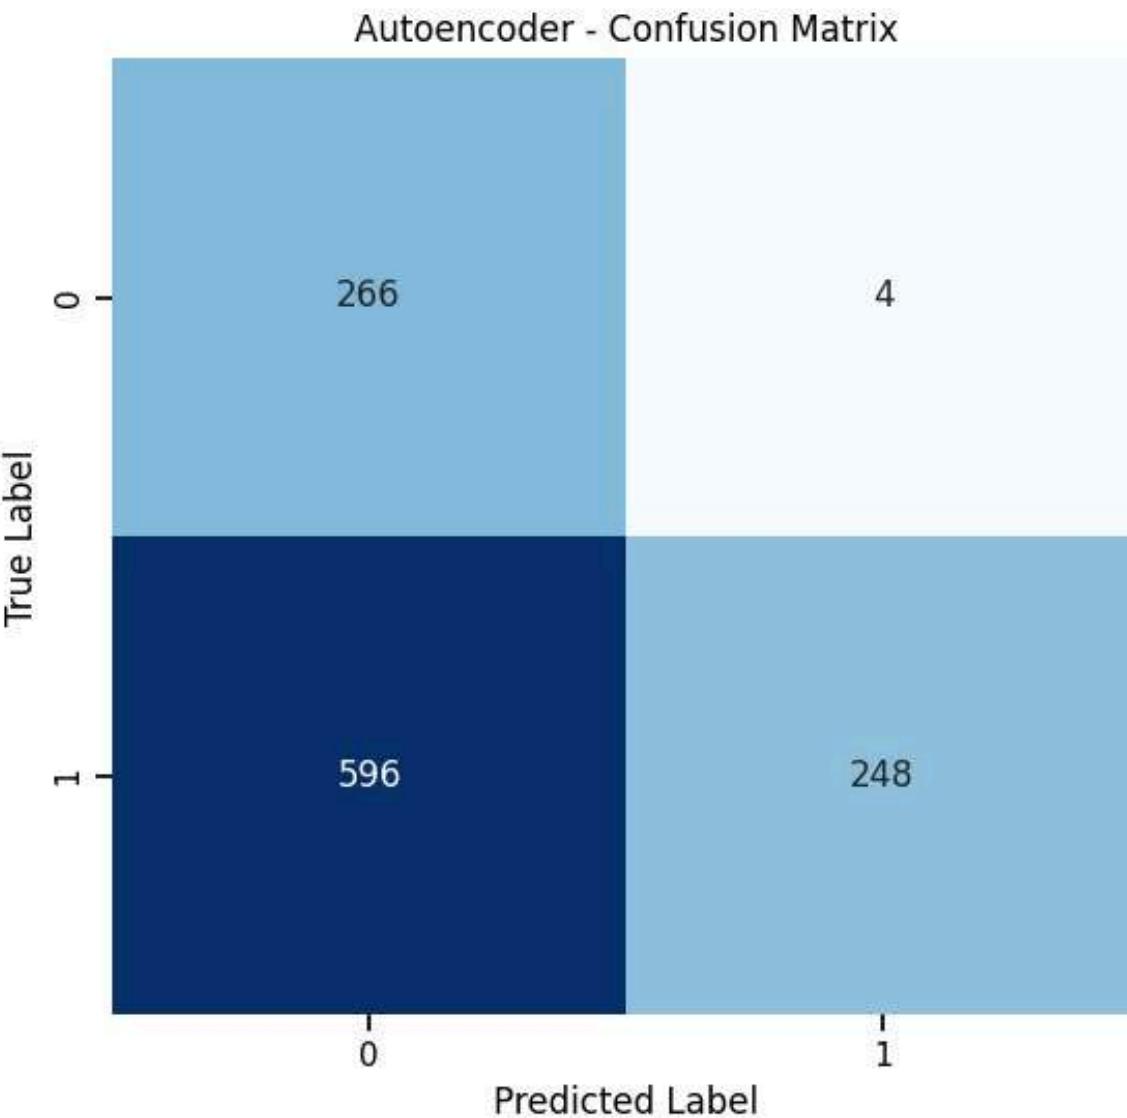
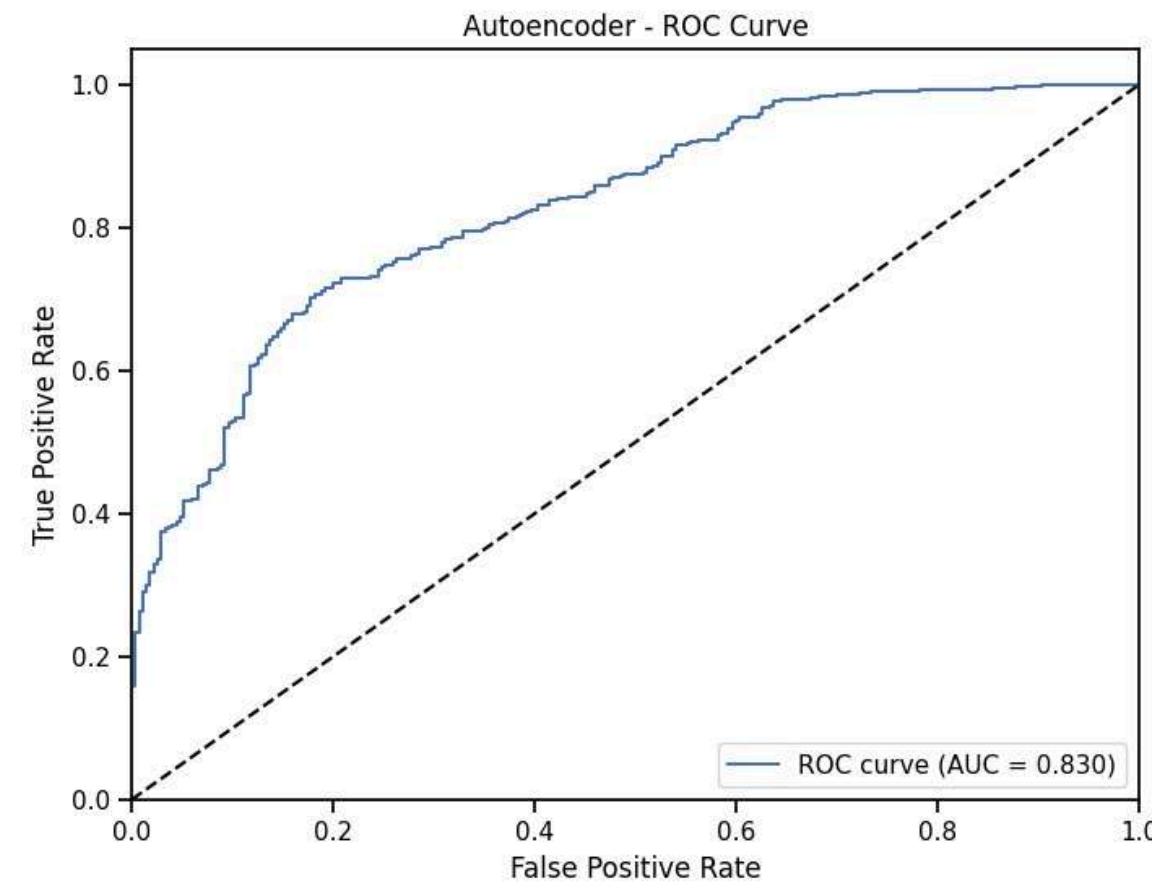
75% of actual anomalies correctly identified

F1 Score

12% false alarms, primarily during transitional periods

Q Anomaly Detection Results

Transaction Spikes Retry Spikes Error Rate Anomalies



Limitations



Data Augmentation

- ✓ Relied solely on oversampling of known “anomalous” windows—no generative noise or SMOTE—so synthetic patterns may lack real-world variability.
- ✓ Augmented data still originates from a single-day snapshot, limiting coverage of seasonal or multi-day dynamics.



Feature Engineering

- ✓ Lack of external enrichments (weather, calendar events) means models miss potentially predictive context.



Forecasting Models

- ✓ Hyperparameters and architectures were fixed (no grid search or cross-validation), so models may not be optimally tuned.



Anomaly Detection

- ✓ 241 anomalies go unnoticed (FN = 241)
- ✓ Trained autoencoder threshold set by a single percentile cut—trade-off between false positives and negatives may not generalize.

Key Takeaways



Threshold tuning or cost-sensitive decision rule



Augment training data with more anomaly examples